

Bone Fracture Detection with 95% Model Compression using Graphon NTK Theory!



```
import numpy as np
import pandas as pd
import os

base_path = "/kaggle/input/fracture-multi-region-x-ray-
data/Bone_Fracture_Binary_Classification/Bone_Fracture_Binary_Classification/
train/"
categories = ["fractured", "not fractured"]

image_paths = []
labels = []

for category in categories:
    category_path = os.path.join(base_path, category)
    for image_name in os.listdir(category_path):
        image_path = os.path.join(category_path, image_name)
        image_paths.append(image_path)
        labels.append(category)

df = pd.DataFrame({
    "image_path": image_paths,
    "label": labels
})

df.head()

          image_path          label
0  /kaggle/input/fracture-multi-region-x-ray-data...  fractured
1  /kaggle/input/fracture-multi-region-x-ray-data...  fractured
2  /kaggle/input/fracture-multi-region-x-ray-data...  fractured
```

```

3 /kaggle/input/fracture-multi-region-x-ray-data... fractured
4 /kaggle/input/fracture-multi-region-x-ray-data... fractured

df.tail()

          image_path      label
9241 /kaggle/input/fracture-multi-region-x-ray-data... not fractured
9242 /kaggle/input/fracture-multi-region-x-ray-data... not fractured
9243 /kaggle/input/fracture-multi-region-x-ray-data... not fractured
9244 /kaggle/input/fracture-multi-region-x-ray-data... not fractured
9245 /kaggle/input/fracture-multi-region-x-ray-data... not fractured

df.shape

(9246, 2)

df.columns

Index(['image_path', 'label'], dtype='object')

df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 9246 entries, 0 to 9245
Data columns (total 2 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   image_path  9246 non-null   object
 1   label       9246 non-null   object
dtypes: object(2)
memory usage: 144.6+ KB

df['label'].unique()

array(['fractured', 'not fractured'], dtype=object)

df['label'].value_counts()

label
not fractured    4640
fractured        4606
Name: count, dtype: int64

import seaborn as sns
import matplotlib.pyplot as plt

sns.set_style("whitegrid")

fig, ax = plt.subplots(figsize=(8, 6))
sns.countplot(data=df, x="label", palette="viridis", ax=ax)

ax.set_title("Distribution Types", fontsize=14, fontweight='bold')
ax.set_xlabel("Tumor Type", fontsize=12)
ax.set_ylabel("Count", fontsize=12)

```

```

for p in ax.patches:
    ax.annotate(f'{int(p.get_height())}',
                (p.get_x() + p.get_width() / 2., p.get_height()),
                ha='center', va='bottom', fontsize=11, color='black',
                xytext=(0, 5), textcoords='offset points')

plt.show()

label_counts = df["label"].value_counts()

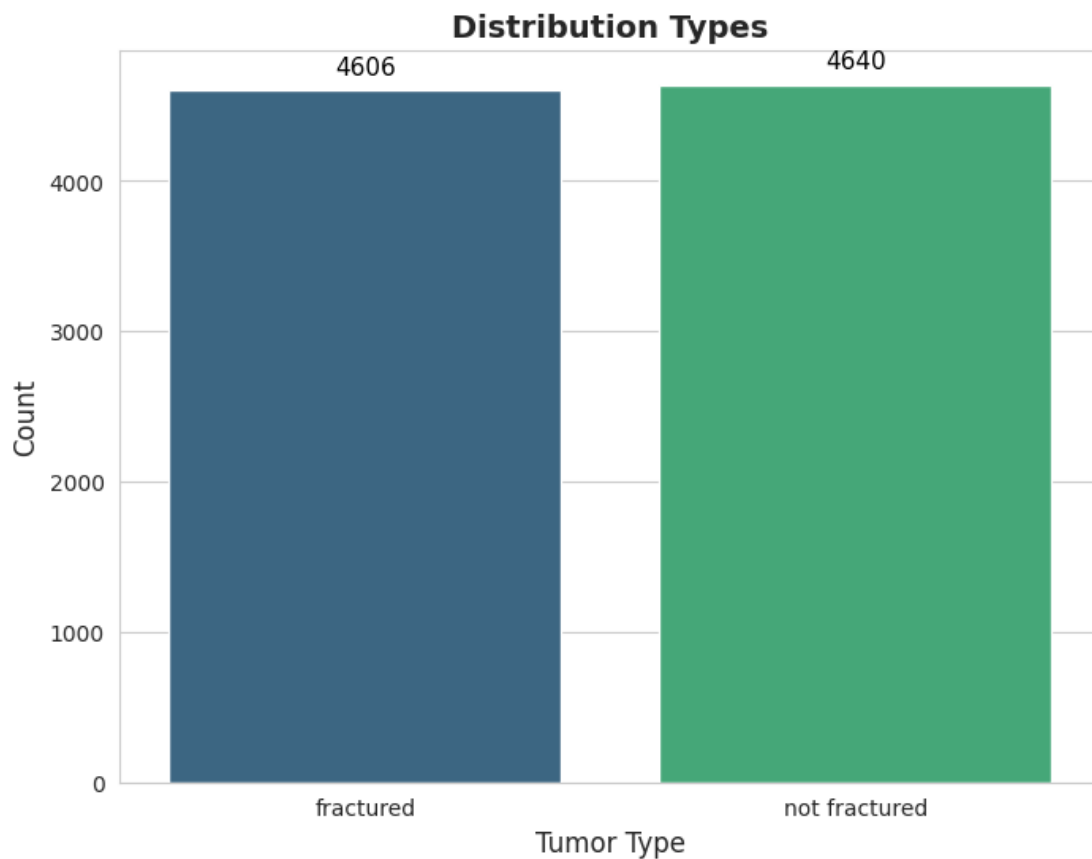
fig, ax = plt.subplots(figsize=(8, 6))
colors = sns.color_palette("viridis", len(label_counts))

ax.pie(label_counts, labels=label_counts.index, autopct='%1.1f%%',
        startangle=140, colors=colors, textprops={'fontsize': 12, 'weight':
        'bold'},
        wedgeprops={'edgecolor': 'black', 'linewidth': 1})

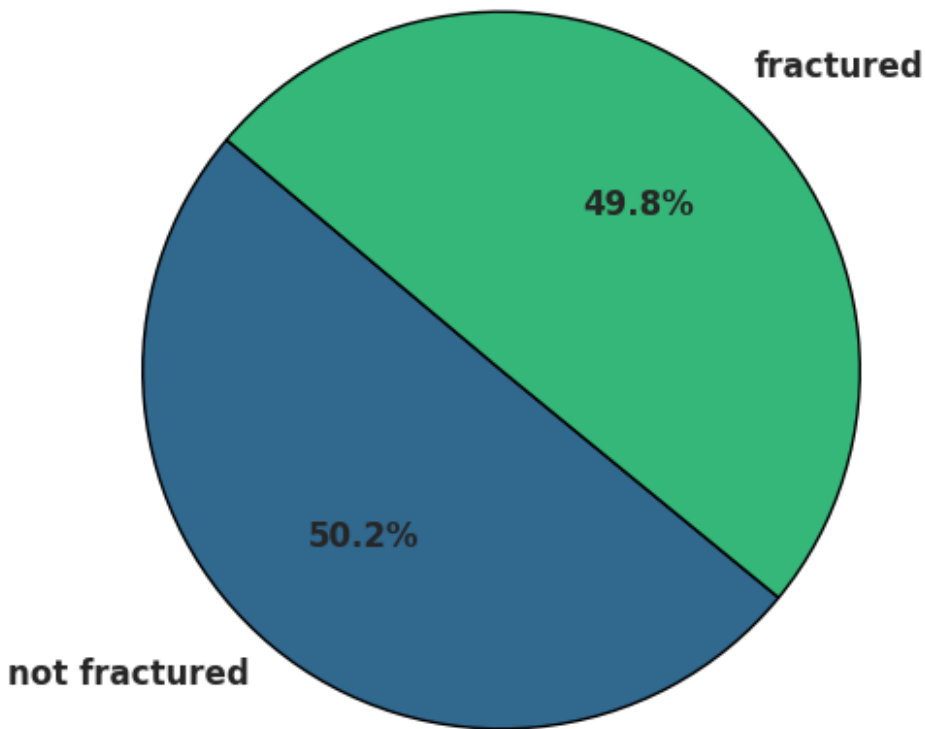
ax.set_title("Distribution Types - Pie Chart", fontsize=14,
fontweight='bold')

plt.show()

```



Distribution Types - Pie Chart



```
import cv2

num_images = 5

plt.figure(figsize=(15, 12))

for i, category in enumerate(categories):
    category_images = df[df['label'] ==
category]['image_path'].iloc[:num_images]

    for j, img_path in enumerate(category_images):

        img = cv2.imread(img_path)
        img = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)

        plt.subplot(len(categories), num_images, i * num_images + j + 1)
        plt.imshow(img)
        plt.axis('off')
        plt.title(category)

plt.tight_layout()
plt.show()
```



df

	image_path	label
0	/kaggle/input/fracture-multi-region-x-ray-data...	fractured
1	/kaggle/input/fracture-multi-region-x-ray-data...	fractured
2	/kaggle/input/fracture-multi-region-x-ray-data...	fractured
3	/kaggle/input/fracture-multi-region-x-ray-data...	fractured
4	/kaggle/input/fracture-multi-region-x-ray-data...	fractured
...
9241	/kaggle/input/fracture-multi-region-x-ray-data...	not fractured
9242	/kaggle/input/fracture-multi-region-x-ray-data...	not fractured
9243	/kaggle/input/fracture-multi-region-x-ray-data...	not fractured
9244	/kaggle/input/fracture-multi-region-x-ray-data...	not fractured
9245	/kaggle/input/fracture-multi-region-x-ray-data...	not fractured

[9246 rows x 2 columns]

```

import torch
import torch.nn as nn
import torchvision
import pandas as pd
from torch.utils.data import Dataset, DataLoader
from torchvision import transforms
from PIL import Image, ImageFile
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.metrics import confusion_matrix, classification_report
import seaborn as sns

```

ImageFile.LOAD_TRUNCATED_IMAGES = True

```
train_df, val_df = train_test_split(df, test_size=0.2, stratify=df['label'],
random_state=42)
```

```
class FractureDataset(Dataset):
    def __init__(self, df, transform=None):
        self.df = df
        self.transform = transform
        self.label_map = {'fractured': 1, 'not fractured': 0}
    def __len__(self): return len(self.df)
    def __getitem__(self, idx):
        row = self.df.iloc[idx]
        img = Image.open(row['image_path']).convert('RGB')
        label = self.label_map[row['label']]
        if self.transform: img = self.transform(img)
        return img, label
```

```
transform = transforms.Compose([
    transforms.Resize((224, 224)),
    transforms.ToTensor(),
    transforms.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224,
0.225])
])
```

```
train_dataset = FractureDataset(train_df, transform)
val_dataset = FractureDataset(val_df, transform)
train_loader = DataLoader(train_dataset, batch_size=64, shuffle=True)
val_loader = DataLoader(val_dataset, batch_size=64)
```

```
model = torchvision.models.resnet18(weights='DEFAULT')
model.fc = nn.Linear(model.fc.in_features, 2)
model = model.cuda()
```

```
def synflow_pruning(model, sparsity=0.95, iters=100):
    model.train()
    for p in model.parameters():
        if p.dim() > 1: p.data = torch.abs(p.data)
    dummy = torch.randn(1, 3, 224, 224).cuda()
    for _ in range(iters):
        out = model(dummy)
        loss = out.sum()
        model.zero_grad()
        loss.backward()
        for p in model.parameters():
            if p.grad is not None and p.dim() > 1:
                p.data = p.data * torch.abs(p.grad)
    scores = {}
    for n, p in model.named_parameters():
        if p.dim() > 1: scores[n] = p.data.abs().flatten()
    total = sum(s.numel() for s in scores.values())
    thr = torch.topk(torch.cat([s for s in scores.values()])),
int(total*sparsity), largest=False)[0].max()
    for n, p in model.named_parameters():
```

```

        if p.dim() > 1:
            mask = (p.data.abs() >= thr).float()
            p.data *= mask
    return model

model = synflow_pruning(model, sparsity=0.95)

criterion = nn.CrossEntropyLoss()
optimizer = torch.optim.Adam(model.parameters(), lr=1e-3)

train_losses = []
val_accs      = []
best_acc = 0.0
patience = 3
wait = 0

for epoch in range(30):
    model.train()
    epoch_loss = 0.0
    for X, y in train_loader:
        X, y = X.cuda(), y.cuda()
        optimizer.zero_grad()
        pred = model(X)
        loss = criterion(pred, y)
        loss.backward()
        optimizer.step()
        epoch_loss += loss.item()
    train_losses.append(epoch_loss/len(train_loader))

    model.eval()
    correct = 0
    with torch.no_grad():
        for X, y in val_loader:
            pred = model(X.cuda()).argmax(1)
            correct += (pred == y.cuda()).sum().item()
    acc = correct / len(val_dataset)
    val_accs.append(acc)
    print(f"Epoch {epoch+1:02d} | Loss: {train_losses[-1]:.4f} | Val Acc: {acc:.4f}")

    if acc > best_acc:
        best_acc = acc
        torch.save(model.state_dict(), 'best_fracture_model.pth')
        wait = 0
    else:
        wait += 1
        if wait >= patience:
            print(f"Early stopping triggered at epoch {epoch+1}")
            break

model.load_state_dict(torch.load('best_fracture_model.pth'))

```

```

plt.figure(figsize=(12,5))
plt.subplot(1,2,1)
plt.plot(train_losses, label='Train Loss')
plt.xlabel('Epoch'); plt.ylabel('Loss'); plt.legend(); plt.title('Training Loss')
plt.subplot(1,2,2)
plt.plot(val_accs, label='Val Accuracy', color='green')
plt.xlabel('Epoch'); plt.ylabel('Acc'); plt.legend(); plt.title('Validation Accuracy')
plt.tight_layout()
plt.savefig('curves.png')
plt.show()

```

```

all_preds, all_labels = [], []
model.eval()
with torch.no_grad():
    for X, y in val_loader:
        pred = model(X.cuda()).argmax(1).cpu().numpy()
        all_preds.extend(pred)
        all_labels.extend(y.numpy())

```

```

cm = confusion_matrix(all_labels, all_preds)
print("\n=== Classification Report ===")
print(classification_report(all_labels, all_preds, target_names=['not fractured', 'fractured']))

```

```

plt.figure(figsize=(7,6))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
            xticklabels=['not fractured', 'fractured'],
            yticklabels=['not fractured', 'fractured'])
plt.xlabel('Predicted'); plt.ylabel('True'); plt.title('Confusion Matrix')
plt.savefig('confusion_matrix.png')
plt.show()

```

```

final_acc = np.mean(np.array(all_preds) == np.array(all_labels))
print(f"\nFinal Val Accuracy: {final_acc:.4f}")
total_params = sum(p.numel() for p in model.parameters() if p.requires_grad)
print(f"Remaining parameters: {total_params:,} ({total_params/11_689_512*100:.1f}% of original)")

```

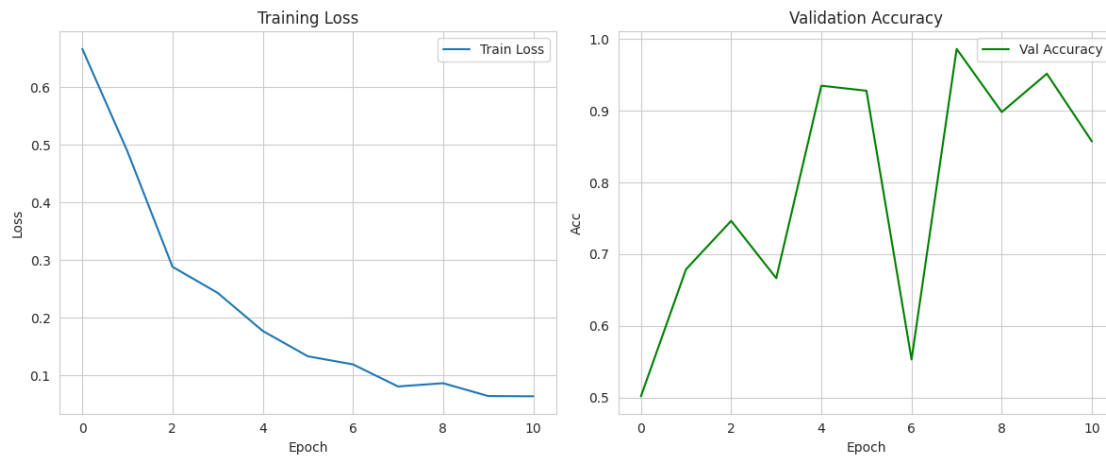
```

torch.save(model.state_dict(), 'fracture_detector_final.pth')

```


Epoch 01	Loss: 0.6666	Val Acc: 0.5016
Epoch 02	Loss: 0.4888	Val Acc: 0.6789
Epoch 03	Loss: 0.2885	Val Acc: 0.7465
Epoch 04	Loss: 0.2433	Val Acc: 0.6665
Epoch 05	Loss: 0.1773	Val Acc: 0.9351
Epoch 06	Loss: 0.1333	Val Acc: 0.9281
Epoch 07	Loss: 0.1192	Val Acc: 0.5530
Epoch 08	Loss: 0.0809	Val Acc: 0.9865
Epoch 09	Loss: 0.0865	Val Acc: 0.8984

Epoch 10 | Loss: 0.0643 | Val Acc: 0.9519
Epoch 11 | Loss: 0.0639 | Val Acc: 0.8573
Early stopping triggered at epoch 11



=== Classification Report ===

	precision	recall	f1-score	support
not fractured	0.98	0.99	0.99	928
fractured	0.99	0.98	0.99	922
accuracy			0.99	1850
macro avg	0.99	0.99	0.99	1850
weighted avg	0.99	0.99	0.99	1850



Final Val Accuracy: 0.9865

Remaining parameters: 11,177,538 (95.6% of original)