


Chest X-Ray Classification using GLCM, Wavelet Transforms, and CLAHE



(a) Normal



(b) Pneumonia



(c) COVID-19

```
import numpy as np
import pandas as pd
import os

import warnings
warnings.filterwarnings('ignore')

root_path = '/kaggle/input/chest-x-ray-images-normal-and-
pneumonia/chest_xray/train/'

image_paths = []
labels = []

for label in os.listdir(root_path):
    label_path = os.path.join(root_path, label)
    if os.path.isdir(label_path):
        for img_file in os.listdir(label_path):
            if img_file.lower().endswith(('.png', '.jpg', '.jpeg')):
                image_paths.append(os.path.join(label_path, img_file))
                labels.append(label)

df = pd.DataFrame({'image_path': image_paths, 'label': labels})

df.head()

          image_path      label
0  /kaggle/input/chest-x-ray-images-normal-and-pn...  PNEUMONIA
1  /kaggle/input/chest-x-ray-images-normal-and-pn...  PNEUMONIA
2  /kaggle/input/chest-x-ray-images-normal-and-pn...  PNEUMONIA
```

```
3 /kaggle/input/chest-x-ray-images-normal-and-pn... PNEUMONIA
4 /kaggle/input/chest-x-ray-images-normal-and-pn... PNEUMONIA
```

```
df.tail()
```

```
          image_path  label
5211 /kaggle/input/chest-x-ray-images-normal-and-pn...  NORMAL
5212 /kaggle/input/chest-x-ray-images-normal-and-pn...  NORMAL
5213 /kaggle/input/chest-x-ray-images-normal-and-pn...  NORMAL
5214 /kaggle/input/chest-x-ray-images-normal-and-pn...  NORMAL
5215 /kaggle/input/chest-x-ray-images-normal-and-pn...  NORMAL
```

```
df.shape
```

```
(5216, 2)
```

```
df.columns
```

```
Index(['image_path', 'label'], dtype='object')
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5216 entries, 0 to 5215
Data columns (total 2 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   image_path  5216 non-null   object
 1   label       5216 non-null   object
dtypes: object(2)
memory usage: 81.6+ KB
```

```
df['label'].unique()
```

```
array(['PNEUMONIA', 'NORMAL'], dtype=object)
```

```
df['label'].value_counts()
```

```
label
PNEUMONIA    3875
NORMAL       1341
Name: count, dtype: int64
```

```
import seaborn as sns
import matplotlib.pyplot as plt
```

```
sns.set_style("whitegrid")
```

```
fig, ax = plt.subplots(figsize=(8, 6))
sns.countplot(data=df, x="label", palette="viridis", ax=ax)
```

```
ax.set_title("Distribution of Tumor Types", fontsize=14, fontweight='bold')
```

```

ax.set_xlabel("Tumor Type", fontsize=12)
ax.set_ylabel("Count", fontsize=12)

for p in ax.patches:
    ax.annotate(f'{int(p.get_height())}',
                (p.get_x() + p.get_width() / 2., p.get_height()),
                ha='center', va='bottom', fontsize=11, color='black',
                xytext=(0, 5), textcoords='offset points')

plt.show()

label_counts = df["label"].value_counts()


fig, ax = plt.subplots(figsize=(10, 8))
colors = sns.color_palette("viridis", len(label_counts))

ax.pie(label_counts, labels=label_counts.index, autopct='%1.1f%%',
        startangle=140, colors=colors, textprops={'fontsize': 8, 'weight':
        'bold'},
        wedgeprops={'edgecolor': 'black', 'linewidth': 1})


ax.set_title("Distribution of Tumor Types - Pie Chart", fontsize=14,
fontweight='bold')

plt.show()

```



Distribution of Tumor Types - Pie Chart



```
from PIL import Image

num_images = 5

unique_labels = df['label'].unique()

plt.figure(figsize=(15, len(unique_labels) * 3))

for row_idx, label in enumerate(unique_labels):

    label_images = df[df['label'] ==
label].head(num_images)['image_path'].tolist()


    for col_idx, img_path in enumerate(label_images):
        plt_idx = row_idx * num_images + col_idx + 1
```

```

plt.subplot(len(unique_labels), num_images, plt_idx)
img = Image.open(img_path)
plt.imshow(img)
plt.axis('off')
if col_idx == 2:
    plt.title(label, fontsize=10)

plt.tight_layout()
plt.show()

```



```

max_samples = df['label'].value_counts().max()

balanced_df = df.groupby('label', group_keys=False).apply(
    lambda x: x.sample(n=max_samples, replace=True, random_state=42)
).reset_index(drop=True)

```

```
balanced_df = balanced_df[['image_path', 'label']]
```

```
df = balanced_df
```

```
df
```

	image_path	label
0	/kaggle/input/chest-x-ray-images-normal-and-pn...	NORMAL
1	/kaggle/input/chest-x-ray-images-normal-and-pn...	NORMAL
2	/kaggle/input/chest-x-ray-images-normal-and-pn...	NORMAL
3	/kaggle/input/chest-x-ray-images-normal-and-pn...	NORMAL
4	/kaggle/input/chest-x-ray-images-normal-and-pn...	NORMAL
...
7745	/kaggle/input/chest-x-ray-images-normal-and-pn...	PNEUMONIA
7746	/kaggle/input/chest-x-ray-images-normal-and-pn...	PNEUMONIA
7747	/kaggle/input/chest-x-ray-images-normal-and-pn...	PNEUMONIA
7748	/kaggle/input/chest-x-ray-images-normal-and-pn...	PNEUMONIA
7749	/kaggle/input/chest-x-ray-images-normal-and-pn...	PNEUMONIA

```
[7750 rows x 2 columns]
```

```

from PIL import Image
import cv2
from skimage.feature import graycomatrix, graycoprops
import pywt
import matplotlib.pyplot as plt

def get_rotation_invariant_glcm_features(image_array, distance=[1],
angles=[0, np.pi/4, np.pi/2, 3*np.pi/4], levels=32):
    image_array = image_array.astype(np.uint8)

    max_val = image_array.max()
    if max_val >= levels:
        quantized_array = np.floor(image_array * ((levels - 1) /
max_val)).astype(np.uint8)
    else:
        quantized_array = image_array

    gcom = graycomatrix(
        quantized_array,
        distances=distance,
        angles=angles,
        levels=levels,
        symmetric=True,
        normed=True
    )

    dissimilarity = np.mean(graycoprops(gcom, 'dissimilarity'))
    correlation = np.mean(graycoprops(gcom, 'correlation'))
    homogeneity = np.mean(graycoprops(gcom, 'homogeneity'))
    energy = np.mean(graycoprops(gcom, 'energy'))

    return [dissimilarity, correlation, homogeneity, energy]

def get_wavelet_features(image_array, wavelet='haar', level=1):
    coeffs = pywt.wavedec2(image_array.astype(float), wavelet, level=level)

    LH, HL, HH = coeffs[1]

    lh_energy = np.sum(LH**2)
    hl_energy = np.sum(HL**2)
    hh_energy = np.sum(HH**2)

    return [lh_energy, hl_energy, hh_energy]

def get_multi_feature_embedding(image_path, resize_to=(128, 128), n_blocks=4,
bit_depth=16):
    if resize_to[0] % n_blocks != 0 or resize_to[1] % n_blocks != 0:
        raise ValueError("resize_to dimensions must be divisible by

```

```

n_blocks.")

block_w = resize_to[0] // n_blocks
block_h = resize_to[1] // n_blocks

try:
    img = Image.open(image_path).convert('L')
    img = img.resize(resize_to)
    img_array = np.array(img).astype(np.uint8)

    clahe = cv2.createCLAHE(clipLimit=2.0, tileGridSize=(8,8))
    clahe_img = clahe.apply(img_array)

    embedding = []

    MAX_DISSIMILARITY = 16.0
    MAX_INTENSITY = 255.0
    MAX_WAVELET_ENERGY = 500000.0

    for i in range(n_blocks):
        for j in range(n_blocks):
            block = clahe_img[i*block_h:(i+1)*block_h,
j*block_w:(j+1)*block_w]

            glcm_features = get_rotation_invariant_glcm_features(block)
            wavelet_features = get_wavelet_features(block)
            mean_intensity = np.mean(block)
            std_intensity = np.std(block)

            features = [
                min(glcm_features[0] / MAX_DISSIMILARITY, 1.0),
                (glcm_features[1] + 1.0) / 2.0,
                glcm_features[2],
                glcm_features[3],
                min(wavelet_features[0] / MAX_WAVELET_ENERGY, 1.0),
                min(wavelet_features[1] / MAX_WAVELET_ENERGY, 1.0),
                min(wavelet_features[2] / MAX_WAVELET_ENERGY, 1.0),
                mean_intensity / MAX_INTENSITY,
                std_intensity / MAX_INTENSITY
            ]

            for feat in features:
                scaled_value = int(feat * (2**bit_depth - 1))
                bin_str = bin(scaled_value)[2:].zfill(bit_depth)
                bin_list = [int(bit) for bit in bin_str]
                embedding.extend(bin_list)

    return embedding
except Exception as e:

```



```

        print(f"Error processing {image_path}: {e}")
        return None

def embedding_to_analog(embedding, bit_depth=16):
    if not embedding: return []
    analog = []
    for i in range(0, len(embedding), bit_depth):
        bin_str = ''.join(str(bit) for bit in embedding[i:i+bit_depth])
        value = int(bin_str, 2)
        analog.append(value)
    return analog

def generate_em_wave(analog_values, feature_weights=None, carrier_freq=5,
sampling_rate=50, duration_per_value=0.1, max_value=2**16 - 1):
    if not analog_values: return np.array([0]), np.array([0])

    total_duration = len(analog_values) * duration_per_value
    t = np.linspace(0, total_duration, int(total_duration * sampling_rate),
endpoint=False)
    wave = np.zeros_like(t)
    samples_per_value = int(duration_per_value * sampling_rate)

    if feature_weights is None:
        feature_weights = np.ones(len(analog_values))

    for i, value in enumerate(analog_values):
        amp = value / max_value
        weighted_amp = amp * feature_weights[i]

        start = i * samples_per_value
        end = min(start + samples_per_value, len(t))

        if start < len(t):
            wave[start:end] = weighted_amp * np.sin(2 * np.pi * carrier_freq
* t[start:end])

    return t, wave

df['multi_feature_embedding'] = df['image_path'].apply(lambda path:
get_multi_feature_embedding(path))

```

```
df
```

	image_path	label \
0	/kaggle/input/chest-x-ray-images-normal-and-pn...	NORMAL
1	/kaggle/input/chest-x-ray-images-normal-and-pn...	NORMAL
2	/kaggle/input/chest-x-ray-images-normal-and-pn...	NORMAL
3	/kaggle/input/chest-x-ray-images-normal-and-pn...	NORMAL
4	/kaggle/input/chest-x-ray-images-normal-and-pn...	NORMAL

```

...
7745 /kaggle/input/chest-x-ray-images-normal-and-pn... PNEUMONIA
7746 /kaggle/input/chest-x-ray-images-normal-and-pn... PNEUMONIA
7747 /kaggle/input/chest-x-ray-images-normal-and-pn... PNEUMONIA
7748 /kaggle/input/chest-x-ray-images-normal-and-pn... PNEUMONIA
7749 /kaggle/input/chest-x-ray-images-normal-and-pn... PNEUMONIA

```

```

                                multi_feature_embedding
0      [0, 0, 0, 1, 0, 1, 1, 0, 0, 1, 0, 0, 0, 0, 0, ...
1      [0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 1, 1, 0, 0, 1, 0, ...
2      [0, 0, 0, 1, 0, 1, 1, 1, 0, 0, 0, 0, 1, 0, 0, 0, ...
3      [0, 0, 0, 1, 1, 1, 0, 0, 0, 0, 0, 1, 1, 0, 1, 1, ...
4      [0, 0, 0, 1, 1, 1, 0, 0, 1, 0, 1, 1, 0, 0, 1, 1, ...
...
7745   [0, 0, 0, 1, 0, 1, 1, 0, 0, 0, 0, 1, 0, 1, 0, 1, 0, ...
7746   [0, 0, 0, 1, 1, 0, 1, 0, 1, 0, 1, 0, 1, 1, 0, 1, 0, ...
7747   [0, 0, 0, 1, 1, 0, 1, 1, 0, 1, 1, 0, 1, 0, 1, 0, 1, ...
7748   [0, 0, 0, 1, 1, 1, 0, 1, 1, 1, 0, 1, 0, 1, 1, 1, 1, ...
7749   [0, 0, 0, 1, 0, 1, 0, 1, 0, 1, 1, 1, 1, 1, 0, 0, 1, ...

```

[7750 rows x 3 columns]

```

def embedding_to_analog(embedding, bit_depth=16):
    if not embedding: return []
    analog = []
    for i in range(0, len(embedding), bit_depth):
        bin_str = ''.join(str(bit) for bit in embedding[i:i+bit_depth])
        value = int(bin_str, 2)
        analog.append(value)
    return analog

def generate_em_wave(analog_values, feature_weights=None, carrier_freq=5,
sampling_rate=50, duration_per_value=0.1, max_value=2**16 - 1):
    if not analog_values: return np.array([0]), np.array([0])

    total_duration = len(analog_values) * duration_per_value
    t = np.linspace(0, total_duration, int(total_duration * sampling_rate),
endpoint=False)
    wave = np.zeros_like(t)
    samples_per_value = int(duration_per_value * sampling_rate)

    if feature_weights is None:
        feature_weights = np.ones(len(analog_values))

    for i, value in enumerate(analog_values):
        amp = value / max_value
        weighted_amp = amp * feature_weights[i]

        start = i * samples_per_value

```

```

        end = min(start + samples_per_value, len(t))

        if start < len(t):
            wave[start:end] = weighted_amp * np.sin(2 * np.pi * carrier_freq
* t[start:end])

    return t, wave

df['multi_analog_values'] =
df['multi_feature_embedding'].apply(embedding_to_analog)

N_PLOTS = 5
fig, axes = plt.subplots(N_PLOTS, 1, figsize=(14, 2 * N_PLOTS))
plt.suptitle('Simulated EM Wave from Multi-Feature Embedding (First 5
Images)', fontsize=14, y=1.02)

FEATURES_TO_PLOT = 50

for i in range(N_PLOTS):
    analog_values = df['multi_analog_values'].iloc[i]
    label = df['label'].iloc[i]

    analog_subsample = analog_values[:FEATURES_TO_PLOT]

    t, wave = generate_em_wave(analog_subsample, carrier_freq=10,
sampling_rate=100)

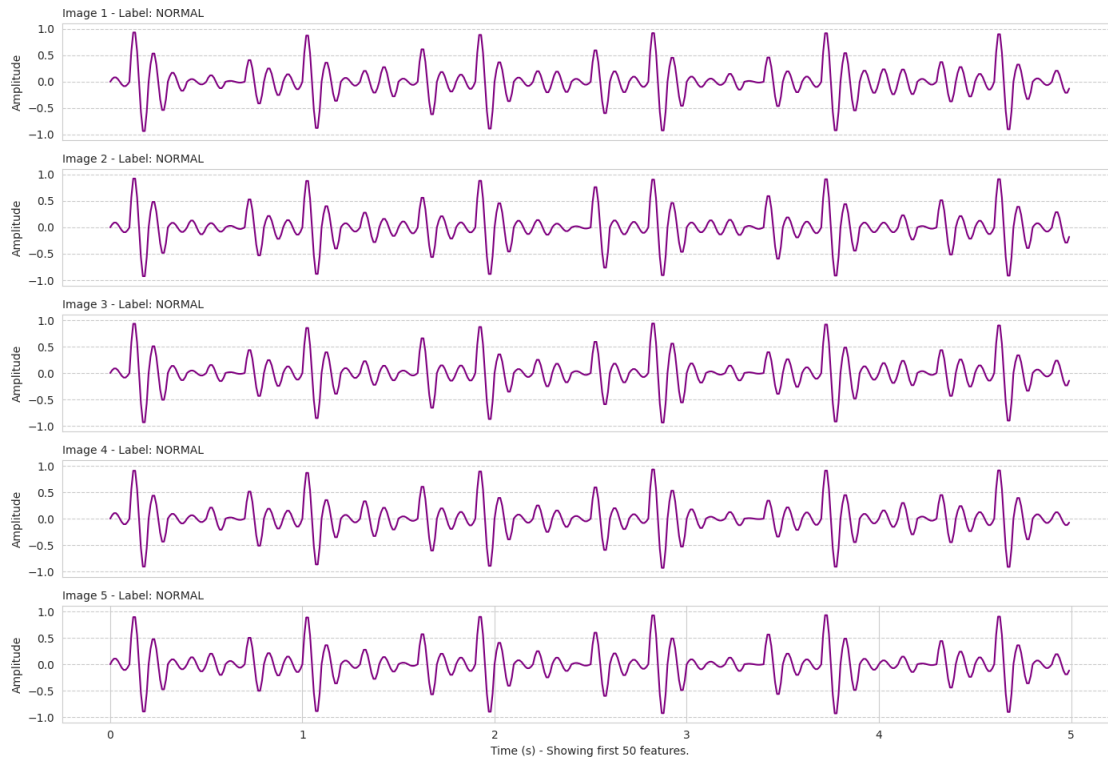
    ax = axes[i]
    ax.plot(t, wave, color='purple')
    ax.set_title(f"Image {i+1} - Label: {label}", fontsize=10, loc='left')
    ax.set_ylabel('Amplitude')
    ax.set_ylim(-1.1, 1.1)
    ax.grid(axis='y', linestyle='--')

    if i < N_PLOTS - 1:
        ax.set_xticks([])

axes[-1].set_xlabel(f'Time (s) - Showing first {FEATURES_TO_PLOT} features.')
plt.tight_layout()
plt.show()

```

Simulated EM Wave from Multi-Feature Embedding (First 5 Images)



```
import matplotlib.pyplot as plt
import matplotlib.gridspec as gridspec

def embedding_to_analog(embedding, bit_depth=16):
    """Converts the binary list back to a list of analog integer values."""
    if not embedding: return []
    analog = []
    for i in range(0, len(embedding), bit_depth):
        bin_str = ''.join(str(bit) for bit in embedding[i:i+bit_depth])
        value = int(bin_str, 2)
        analog.append(value)
    return analog

FEATURE_NAMES = [
    'GLCM: Dissimilarity (Avg)',
    'GLCM: Correlation (Avg)',
    'GLCM: Homogeneity (Avg)',
    'GLCM: Energy (Avg)',
    'Wavelet: LH Energy',
    'Wavelet: HL Energy',
    'Wavelet: HH Energy',
    'Intensity: Mean',
    'Intensity: Std Dev'
]
N_FEATURES_PER_BLOCK = 9
```

```

N_BLOCKS = 4

if 'multi_analog_values' not in df.columns:
    df['multi_analog_values'] =
df['multi_feature_embedding'].apply(embedding_to_analog)

analog_values = df['multi_analog_values'].iloc[0]
image_label = df['label'].iloc[0]

analog_matrix = np.array(analog_values).reshape(N_BLOCKS * N_BLOCKS,
N_FEATURES_PER_BLOCK)

fig = plt.figure(figsize=(12, 12))
gs = gridspec.GridSpec(3, 3, figure=fig, hspace=0.3, wspace=0.3)
plt.suptitle(f'Feature Maps for First Image (Label: {image_label})',
fontsize=16, y=0.95)

for f_idx, feature_name in enumerate(FEATURE_NAMES):
    feature_magnitudes = analog_matrix[:, f_idx]

    feature_map = feature_magnitudes.reshape(N_BLOCKS, N_BLOCKS)

    ax = fig.add_subplot(gs[f_idx // 3, f_idx % 3])

    im = ax.imshow(feature_map, cmap='magma', interpolation='nearest')


    ax.set_title(feature_name, fontsize=10)
    ax.set_xticks([])
    ax.set_yticks([])

    cbar = fig.colorbar(im, ax=ax, fraction=0.045, pad=0.04)
    cbar.ax.tick_params(labelsize=8)

plt.tight_layout(rect=[0, 0, 1, 0.9])
plt.show()

```

Feature Maps for First Image (Label: NORMAL)



df

	image_path	label
0	/kaggle/input/chest-x-ray-images-normal-and-pn...	NORMAL
1	/kaggle/input/chest-x-ray-images-normal-and-pn...	NORMAL
2	/kaggle/input/chest-x-ray-images-normal-and-pn...	NORMAL
3	/kaggle/input/chest-x-ray-images-normal-and-pn...	NORMAL
4	/kaggle/input/chest-x-ray-images-normal-and-pn...	NORMAL
...
7745	/kaggle/input/chest-x-ray-images-normal-and-pn...	PNEUMONIA
7746	/kaggle/input/chest-x-ray-images-normal-and-pn...	PNEUMONIA
7747	/kaggle/input/chest-x-ray-images-normal-and-pn...	PNEUMONIA
7748	/kaggle/input/chest-x-ray-images-normal-and-pn...	PNEUMONIA
7749	/kaggle/input/chest-x-ray-images-normal-and-pn...	PNEUMONIA

multi_feature_embedding \

```

0      [0, 0, 0, 1, 0, 1, 1, 0, 0, 1, 0, 0, 0, 0, 0, ...
1      [0, 0, 0, 1, 1, 0, 0, 0, 0, 1, 1, 0, 0, 1, 0, ...
2      [0, 0, 0, 1, 0, 1, 1, 1, 0, 0, 0, 1, 0, 0, 0, ...
3      [0, 0, 0, 1, 1, 1, 0, 0, 0, 0, 1, 1, 0, 1, 1, ...
4      [0, 0, 0, 1, 1, 1, 0, 0, 1, 0, 1, 1, 0, 0, 1, ...
...
7745   [0, 0, 0, 1, 0, 1, 1, 0, 0, 0, 0, 1, 0, 1, 0, ...
7746   [0, 0, 0, 1, 1, 0, 1, 0, 1, 0, 1, 0, 1, 1, 0, ...
7747   [0, 0, 0, 1, 1, 0, 1, 1, 0, 1, 1, 0, 1, 0, 1, ...
7748   [0, 0, 0, 1, 1, 1, 0, 1, 1, 1, 0, 1, 0, 1, 1, ...
7749   [0, 0, 0, 1, 0, 1, 0, 1, 0, 1, 1, 1, 1, 0, 0, ...

```

```

                                multi_analog_values
0      [5696, 64320, 36841, 11664, 3428, 8345, 1094, ...
1      [6244, 63693, 33233, 5876, 8908, 5479, 2024, 3...
2      [5905, 64272, 34866, 9217, 3188, 10463, 1082, ...
3      [7223, 62548, 29962, 6183, 4374, 14396, 1542, ...
4      [7347, 61600, 32665, 6383, 9102, 13798, 4268, ...
...
7745   [5653, 63222, 35038, 6911, 4154, 6083, 1326, 4...
7746   [6828, 57273, 35833, 7872, 8856, 21497, 4005, ...
7747   [7018, 61047, 32561, 7962, 4034, 13563, 1470, ...
7748   [7639, 61642, 31029, 5852, 5725, 11356, 2023, ...
7749   [5496, 63608, 40484, 11421, 2583, 28977, 2680,...

```

[7750 rows x 4 columns]

```

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense
from tensorflow.keras.optimizers import Adam
from sklearn.metrics import confusion_matrix, classification_report

```

```

X = np.array(df['multi_feature_embedding'].tolist())
le = LabelEncoder()
y = le.fit_transform(df['label'])
target_names = le.classes_

```

```

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
random_state=42)

```

```

model = Sequential([
    Dense(128, activation='relu', input_shape=(X_train.shape[1],)),
    Dense(64, activation='relu'),
    Dense(1, activation='sigmoid')
])

```

```

model.compile(optimizer=Adam(learning_rate=0.001),

```

```

        loss='binary_crossentropy',
        metrics=['accuracy'])

EPOCHS = 5
history = model.fit(
    X_train, y_train,
    epochs=EPOCHS,
    batch_size=32,
    validation_split=0.1,
    verbose=1
)

y_pred_proba = model.predict(X_test)
y_pred = (y_pred_proba > 0.5).astype(int)

plt.figure(figsize=(12, 4))
plt.subplot(1, 2, 1)
plt.plot(history.history['accuracy'], label='Training Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
plt.title('Training and Validation Accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend()
plt.subplot(1, 2, 2)
plt.plot(history.history['loss'], label='Training Loss')
plt.plot(history.history['val_loss'], label='Validation Loss')
plt.title('Training and Validation Loss')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.legend()
plt.show()

cm = confusion_matrix(y_test, y_pred)
plt.figure(figsize=(6, 5))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
            xticklabels=target_names, yticklabels=target_names)
plt.title('Confusion Matrix')
plt.xlabel('Predicted Label')
plt.ylabel('True Label')
plt.show()

report = classification_report(y_test, y_pred, target_names=target_names)
print("\n--- Classification Report ---")
print(report)

2025-10-21 07:25:54.852301: E
external/local_xla/xla/stream_executor/cuda/cuda_fft.cc:477] Unable to
register cuFFT factory: Attempting to register factory for plugin cuFFT when
one has already been registered
WARNING: All log messages before absl::InitializeLog() is called are written

```


to STDERR

```
E0000 00:00:1761031555.257599      37 cuda_dnn.cc:8310] Unable to register
cuDNN factory: Attempting to register factory for plugin cuDNN when one has
already been registered
E0000 00:00:1761031555.392853      37 cuda_blas.cc:1418] Unable to register
cuBLAS factory: Attempting to register factory for plugin cuBLAS when one has
already been registered
I0000 00:00:1761031574.373998      37 gpu_device.cc:2022] Created device
/job:localhost/replica:0/task:0/device:GPU:0 with 13942 MB memory: ->
device: 0, name: Tesla T4, pci bus id: 0000:00:04.0, compute capability: 7.5
I0000 00:00:1761031574.374819      37 gpu_device.cc:2022] Created device
/job:localhost/replica:0/task:0/device:GPU:1 with 13942 MB memory: ->
device: 1, name: Tesla T4, pci bus id: 0000:00:05.0, compute capability: 7.5
```

Epoch 1/5

WARNING: All log messages before absl::InitializeLog() is called are written to STDERR

```
I0000 00:00:1761031578.216555      119 service.cc:148] XLA service
0x7bbd3000d430 initialized for platform CUDA (this does not guarantee that
XLA will be used). Devices:
I0000 00:00:1761031578.217756      119 service.cc:156]   StreamExecutor device
(0): Tesla T4, Compute Capability 7.5
I0000 00:00:1761031578.217776      119 service.cc:156]   StreamExecutor device
(1): Tesla T4, Compute Capability 7.5
I0000 00:00:1761031578.544507      119 cuda_dnn.cc:529] Loaded cuDNN version
90300
```

69/175 ————— 0s 2ms/step - accuracy: 0.6833 - loss: 0.6473

```
I0000 00:00:1761031579.531396      119 device_compiler.h:188] Compiled cluster
using XLA! This line is logged at most once for the lifetime of the process.
```

175/175 ————— 5s 10ms/step - accuracy: 0.7839 - loss: 0.4671 -
val_accuracy: 0.9145 - val_loss: 0.1869

Epoch 2/5

175/175 ————— 1s 3ms/step - accuracy: 0.9700 - loss: 0.0995 -
val_accuracy: 0.9629 - val_loss: 0.1069

Epoch 3/5

175/175 ————— 0s 3ms/step - accuracy: 0.9877 - loss: 0.0438 -
val_accuracy: 0.9806 - val_loss: 0.0814




Epoch 4/5

175/175 ————— 0s 3ms/step - accuracy: 0.9955 - loss: 0.0191 -
val_accuracy: 0.9774 - val_loss: 0.0917

Epoch 5/5

175/175 ————— 1s 3ms/step - accuracy: 0.9997 - loss: 0.0049 -
val_accuracy: 0.9790 - val_loss: 0.0934

49/49 ————— 1s 7ms/step



```

--- Classification Report ---
              precision    recall  f1-score   support

   NORMAL       0.97       0.99       0.98         779
  PNEUMONIA     0.99       0.97       0.98         771

   accuracy              0.98         1550
  macro avg              0.98       0.98       0.98         1550
  
```

weighted avg	0.98	0.98	0.98	1550
--------------	------	------	------	------

```
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from xgboost import XGBClassifier
import time

models = {
    "Logistic Regression (LR)": LogisticRegression(max_iter=1000,
    solver='liblinear', random_state=42),
    "Decision Tree (DT)": DecisionTreeClassifier(max_depth=10,
    random_state=42),
    "Random Forest (RF)": RandomForestClassifier(n_estimators=100,
    max_depth=10, random_state=42, n_jobs=-1),
    "XGBoost (XGB)": XGBClassifier(n_estimators=100, use_label_encoder=False,
    eval_metric='logloss', random_state=42, n_jobs=-1)
}

results = {}

print("--- Starting Model Training and Evaluation ---")
for name, model in models.items():
    start_time = time.time()

    model.fit(X_train, y_train)

    y_pred = model.predict(X_test)

    end_time = time.time()

    results[name] = {
        'model': model,
        'predictions': y_pred,
        'report': classification_report(y_test, y_pred,
    target_names=target_names, output_dict=True),
        'time': end_time - start_time
    }

    print(f"\n{name} Trained in: {results[name]['time']:.2f} seconds")

fig, axes = plt.subplots(2, 2, figsize=(14, 12))
axes = axes.flatten()
plt.suptitle('Model Comparison: Confusion Matrices', fontsize=16)

summary_data = []
for i, (name, res) in enumerate(results.items()):
    y_pred = res['predictions']
```

```

cm = confusion_matrix(y_test, y_pred)
acc = res['report']['accuracy']

sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
            xticklabels=target_names, yticklabels=target_names,
ax=axes[i])
axes[i].set_title(f"{name}\nAccuracy: {acc:.4f}", fontsize=12)
axes[i].set_xlabel('Predicted Label')
axes[i].set_ylabel('True Label')

summary_data.append([
    name,
    f"{acc:.4f}",
    f"{res['report']['PNEUMONIA']['f1-score']:.4f}",
    f"{res['report']['NORMAL']['recall']:.4f}",
    f"{res['report']['PNEUMONIA']['recall']:.4f}",
    f"{res['time']:.2f}s"
])

plt.tight_layout(rect=[0, 0.03, 1, 0.95])
plt.show()

print("\n" + "="*80)
print("                DETAILED CLASSIFICATION REPORTS (Macro Avg F1-Score)")
print("="*80)
for name, res in results.items():
    print(f"\n--- {name} ---")
    print(classification_report(y_test, res['predictions'],
target_names=target_names))

summary_df = pd.DataFrame(summary_data, columns=['Model', 'Accuracy',
'Pneumonia F1', 'Normal Recall', 'Pneumonia Recall', 'Train Time'])
print("\n" + "="*80)
print("                SUMMARY PERFORMANCE TABLE")
print("="*80)
print(summary_df.to_string(index=False))

--- Starting Model Training and Evaluation ---

Logistic Regression (LR) Trained in: 1.85 seconds


Decision Tree (DT) Trained in: 1.37 seconds

Random Forest (RF) Trained in: 1.09 seconds

XGBoost (XGB) Trained in: 2.31 seconds

```

Model Comparison: Confusion Matrices



DETAILED CLASSIFICATION REPORTS (Macro Avg F1-Score)

```
--- Logistic Regression (LR) ---
              precision    recall  f1-score   support

   NORMAL         0.96         0.99         0.98         779
  PNEUMONIA         0.99         0.96         0.98         771

   accuracy              0.98
  macro avg              0.98
 weighted avg              0.98
```

```
--- Decision Tree (DT) ---
              precision    recall  f1-score   support
```

NORMAL	0.93	0.98	0.95	779
PNEUMONIA	0.98	0.92	0.95	771
accuracy			0.95	1550
macro avg	0.95	0.95	0.95	1550
weighted avg	0.95	0.95	0.95	1550

--- Random Forest (RF) ---

	precision	recall	f1-score	support
NORMAL	0.93	0.99	0.96	779
PNEUMONIA	0.99	0.92	0.96	771
accuracy			0.96	1550
macro avg	0.96	0.96	0.96	1550
weighted avg	0.96	0.96	0.96	1550

--- XGBoost (XGB) ---

	precision	recall	f1-score	support
NORMAL	0.96	0.99	0.98	779
PNEUMONIA	0.99	0.96	0.98	771
accuracy			0.98	1550
macro avg	0.98	0.98	0.98	1550
weighted avg	0.98	0.98	0.98	1550

=====

===

SUMMARY PERFORMANCE TABLE

=====

===

	Model	Accuracy	Pneumonia F1	Normal Recall	Pneumonia Recall
Train Time					
Logistic Regression (LR)	0.9768	0.9762	0.9949	0.9585	
1.85s					
Decision Tree (DT)	0.9529	0.9511	0.9846	0.9209	
1.37s					
Random Forest (RF)	0.9587	0.9569	0.9949	0.9222	
1.09s					
XGBoost (XGB)	0.9781	0.9776	0.9936	0.9624	
2.31s					